

Car_Price_Assignment

March 13, 2025

```
[278]: #Import the libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[280]: # Load Data
df = pd.read_csv('CarPrice_Assignment.csv')
df
```

```
[280]:
```

	car_ID	symboling	CarName	fueltype	aspiration	\
0	1	3	alfa-romero giulia	gas	std	
1	2	3	alfa-romero stelvio	gas	std	
2	3	1	alfa-romero Quadrifoglio	gas	std	
3	4	2	audi 100 ls	gas	std	
4	5	2	audi 100ls	gas	std	
..	
200	201	-1	volvo 145e (sw)	gas	std	
201	202	-1	volvo 144ea	gas	turbo	
202	203	-1	volvo 244dl	gas	std	
203	204	-1	volvo 246	diesel	turbo	
204	205	-1	volvo 264gl	gas	turbo	

	doornumber	carbody	drivewheel	engine	location	wheelbase	...	\
0	two	convertible	rwd	front	88.6	...		
1	two	convertible	rwd	front	88.6	...		
2	two	hatchback	rwd	front	94.5	...		
3	four	sedan	fwd	front	99.8	...		
4	four	sedan	4wd	front	99.4	...		
..		
200	four	sedan	rwd	front	109.1	...		
201	four	sedan	rwd	front	109.1	...		
202	four	sedan	rwd	front	109.1	...		
203	four	sedan	rwd	front	109.1	...		
204	four	sedan	rwd	front	109.1	...		

	engine	size	fuel	system	boreratio	stroke	compressionratio	horsepower	\
0	130	mpfi	3.47	2.68	9.0	111			

1	130	mpfi	3.47	2.68	9.0	111
2	152	mpfi	2.68	3.47	9.0	154
3	109	mpfi	3.19	3.40	10.0	102
4	136	mpfi	3.19	3.40	8.0	115
..
200	141	mpfi	3.78	3.15	9.5	114
201	141	mpfi	3.78	3.15	8.7	160
202	173	mpfi	3.58	2.87	8.8	134
203	145	idi	3.01	3.40	23.0	106
204	141	mpfi	3.78	3.15	9.5	114

	peakrpm	citympg	highwaympg	price
0	5000	21	27	13495.0
1	5000	21	27	16500.0
2	5000	19	26	16500.0
3	5500	24	30	13950.0
4	5500	18	22	17450.0
..
200	5400	23	28	16845.0
201	5300	19	25	19045.0
202	5500	18	23	21485.0
203	4800	26	27	22470.0
204	5400	19	25	22625.0

[205 rows x 26 columns]

```
[282]: #Inspecting the data frame.
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null   int64
1   symboling              205 non-null   int64
2   CarName               205 non-null   object
3   fueltype              205 non-null   object
4   aspiration            205 non-null   object
5   doornumber            205 non-null   object
6   carbody               205 non-null   object
7   drivewheel           205 non-null   object
8   enginelocation        205 non-null   object
9   wheelbase             205 non-null   float64
10  carlength             205 non-null   float64
11  carwidth              205 non-null   float64
12  carheight             205 non-null   float64
13  curbweight            205 non-null   int64
```

```

14  enginetype      205 non-null  object
15  cylindernumber  205 non-null  object
16  enginesize      205 non-null  int64
17  fuelsystem      205 non-null  object
18  boreratio       205 non-null  float64
19  stroke          205 non-null  float64
20  compressionratio 205 non-null  float64
21  horsepower      205 non-null  int64
22  peakrpm         205 non-null  int64
23  citympg         205 non-null  int64
24  highwaympg      205 non-null  int64
25  price           205 non-null  float64

```

dtypes: float64(8), int64(8), object(10)

memory usage: 41.8+ KB

None

```
[284]: print(df.describe())
```

	car_ID	symboling	wheelbase	carlength	carwidth	carheight	\
count	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	
mean	103.000000	0.834146	98.756585	174.049268	65.907805	53.724878	
std	59.322565	1.245307	6.021776	12.337289	2.145204	2.443522	
min	1.000000	-2.000000	86.600000	141.100000	60.300000	47.800000	
25%	52.000000	0.000000	94.500000	166.300000	64.100000	52.000000	
50%	103.000000	1.000000	97.000000	173.200000	65.500000	54.100000	
75%	154.000000	2.000000	102.400000	183.100000	66.900000	55.500000	
max	205.000000	3.000000	120.900000	208.100000	72.300000	59.800000	

	curbweight	enginesize	boreratio	stroke	compressionratio	\
count	205.000000	205.000000	205.000000	205.000000	205.000000	
mean	2555.565854	126.907317	3.329756	3.255415	10.142537	
std	520.680204	41.642693	0.270844	0.313597	3.972040	
min	1488.000000	61.000000	2.540000	2.070000	7.000000	
25%	2145.000000	97.000000	3.150000	3.110000	8.600000	
50%	2414.000000	120.000000	3.310000	3.290000	9.000000	
75%	2935.000000	141.000000	3.580000	3.410000	9.400000	
max	4066.000000	326.000000	3.940000	4.170000	23.000000	

	horsepower	peakrpm	citympg	highwaympg	price
count	205.000000	205.000000	205.000000	205.000000	205.000000
mean	104.117073	5125.121951	25.219512	30.751220	13276.710571
std	39.544167	476.985643	6.542142	6.886443	7988.852332
min	48.000000	4150.000000	13.000000	16.000000	5118.000000
25%	70.000000	4800.000000	19.000000	25.000000	7788.000000
50%	95.000000	5200.000000	24.000000	30.000000	10295.000000
75%	116.000000	5500.000000	30.000000	34.000000	16503.000000
max	288.000000	6600.000000	49.000000	54.000000	45400.000000

```
[286]: print(df.head())
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	\
0	1	3	alfa-romero giulia	gas	std	two	
1	2	3	alfa-romero stelvio	gas	std	two	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	
3	4	2	audi 100 ls	gas	std	four	
4	5	2	audi 100ls	gas	std	four	

	carbody	drivewheel	engine location	wheelbase	...	enginesize	\
0	convertible	rwd	front	88.6	...	130	
1	convertible	rwd	front	88.6	...	130	
2	hatchback	rwd	front	94.5	...	152	
3	sedan	fwd	front	99.8	...	109	
4	sedan	4wd	front	99.4	...	136	

	fuelsystem	bore ratio	stroke	compression ratio	horsepower	peakrpm	citympg	\
0	mpfi	3.47	2.68	9.0	111	5000	21	
1	mpfi	3.47	2.68	9.0	111	5000	21	
2	mpfi	2.68	3.47	9.0	154	5000	19	
3	mpfi	3.19	3.40	10.0	102	5500	24	
4	mpfi	3.19	3.40	8.0	115	5500	18	

	highwaympg	price
0	27	13495.0
1	27	16500.0
2	26	16500.0
3	30	13950.0
4	22	17450.0

[5 rows x 26 columns]

```
[288]: #Data Cleaning
```

```
def check_null_percentage(df):
    null_counts = df.isnull().sum()
    null_percentage = (null_counts / len(df)) * 100
    return pd.DataFrame({'Null Count': null_counts, 'Null Percentage':
↪null_percentage})
```

```
[290]: null_df = check_null_percentage(df)
print(null_df)
```

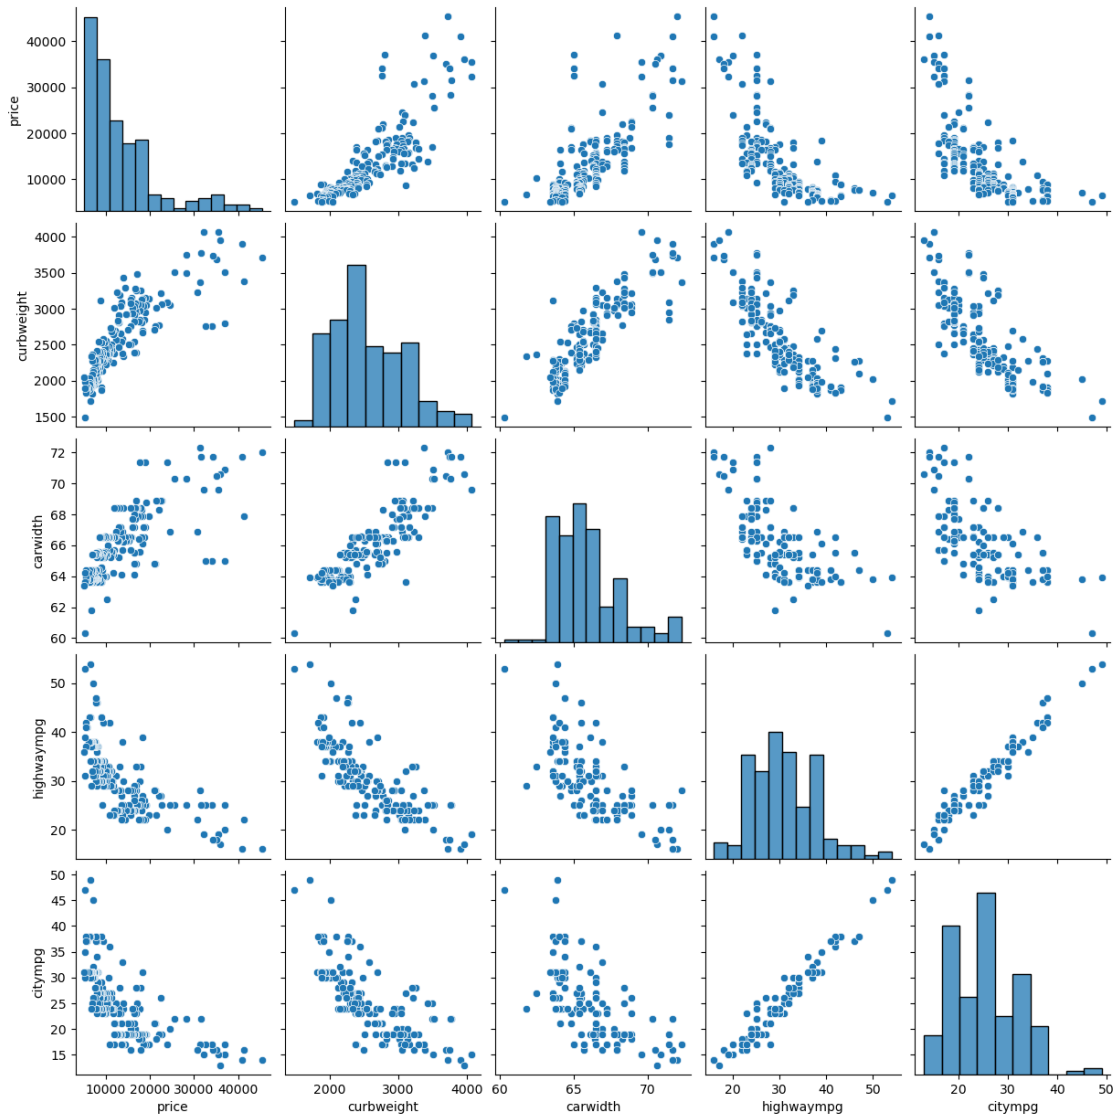
	Null Count	Null Percentage
car_ID	0	0.0
symboling	0	0.0
CarName	0	0.0
fueltype	0	0.0

aspiration	0	0.0
doornumber	0	0.0
carbody	0	0.0
drivewheel	0	0.0
engine location	0	0.0
wheelbase	0	0.0
carlength	0	0.0
carwidth	0	0.0
carheight	0	0.0
curbweight	0	0.0
enginetype	0	0.0
cylindernumber	0	0.0
enginesize	0	0.0
fuelsystem	0	0.0
boreratio	0	0.0
stroke	0	0.0
compressionratio	0	0.0
horsepower	0	0.0
peakrpm	0	0.0
citympg	0	0.0
highwaympg	0	0.0
price	0	0.0


```
[292]: # Dropping unwanted columns
drop_columns = ['car_ID'] # Unnecessary column
df = df.drop(columns=drop_columns, axis=1)
```

```
[294]: #Sorting Column
df = df.sort_values(by='price', ascending=False)
```

```
[296]: # Data Visualization
sns.pairplot(df[['price', 'curbweight', 'carwidth', 'highwaympg', 'citympg']])
plt.show()
```



```
[297]: plt.figure(figsize=(10, 6))
numeric_df = df.select_dtypes(include=[np.number])
sns.heatmap(numeric_df.corr(), annot=True, cmap='coolwarm')
plt.show()
```



```
[298]: #Data Preparation
# Fetching the all company names

car_names = df['CarName'].unique().tolist()
print(car_names)
```

```
['buick regal sport coupe (turbo)', 'bmw x5', 'buick century special', 'porsche
boxter', 'bmw x3', 'jaguar xk', 'jaguar xf', 'buick skylark', 'buick opel isuzu
deluxe', 'porsche cayenne', 'porsche panamera', 'jaguar xj', 'buick skyhawk',
'bmw x4', 'buick century luxury (sw)', 'buick century', 'buick electra 225
custom', 'bmw z4', 'audi 4000', 'volvo 264gl', 'volvo 246', 'porsche macan',
'volvo 244dl', 'bmw x1', 'nissan kicks', 'volvo 144ea', 'volvo diesel', 'audi
5000', 'saab 99e', 'nissan clipper', 'mazda rx-7 gs', 'mazda glc', 'saab 99gle',
'peugeot 604sl', 'peugeot 504', 'audi 5000s (diesel)', 'audi 100ls', 'toyota
cressida', 'nissan teana', 'peugeot 505s turbo diesel', 'bmw 320i', 'volvo 145e
(sw)', 'toyota corolla liftback', 'volvo 245', 'mercury cougar', 'alfa-romero
stelvio', 'alfa-romero Quadrifoglio', 'toyota corona', 'toyota tercel', 'toyota
starlet', 'mazda glc 4', 'audi fox', 'saab 99le', 'mitsubishi g4', 'mitsubishi
mirage g4', 'nissan fuga', 'audi 100 ls', 'volkswagen rabbit custom', 'nissan
otti', 'nissan dayz', 'alfa-romero giulia', 'volkswagen rabbit', 'peugeot 304',
'dodge coronet custom (sw)', 'honda civic', 'plymouth duster', 'mitsubishi
outlander', 'peugeot 504 (sw)', 'volkswagen dasher', 'subaru dl', 'vw dasher',
'toyota corolla', 'subaru r2', 'toyota mark ii', 'mazda 626', 'toyota tercel',
```

```
'isuzu D-Max', 'mazda glc deluxe', 'toyota celica gt', 'mazda glc custom',
'honda civic (auto)', 'honda accord', 'mazda rx-4', 'subaru tribeca', 'vw
rabbit', 'subaru baja', 'renault 5 gtl', 'nissan nv200', 'toyota corolla
tercel', 'volkswagen super beetle', 'toyota celica gt liftback', 'renault 12tl',
'subaru r1', 'honda civic 1300', 'nissan rogue', 'plymouth valiant', 'dodge dart
custom', 'isuzu D-Max V-Cross', 'honda prelude', 'toyota carina', 'dodge coronet
custom', 'volkswagen 411 (sw)', 'mazda glc custom l', 'toyota corona liftback',
'volkswagen type 3', 'mitsubishi pajero', 'nissan note', 'volkswagen model 111',
'volkswagen 1131 deluxe sedan', 'dodge d200', 'plymouth cricket', 'toyota
corolla 1600 (sw)', 'nissan juke', 'subaru brz', 'volkswagen rabbit', 'plymouth
fury gran sedan', 'dodge colt (sw)', 'nissan latio', 'subaru trezia', 'nissan
titan', 'nissan leaf', 'honda civic 1500 gl', 'honda accord lx', 'toyota corolla
1200', 'honda civic cvcc', 'subaru', 'nissan gt-r', 'mitsubishi montero',
'toyota corona hardtop', 'mazda rx2 coupe', 'isuzu MU-X', 'plymouth satellite
custom (sw)', 'dodge colt hardtop', 'chevrolet vega 2300', 'honda accord cvcc',
'dodge challenger se', 'chevrolet monte carlo', 'plymouth fury iii', 'dodge
monaco (sw)', 'mitsubishi lancer', 'mazda glc deluxe', 'dodge rampage', 'Nissan
versa', 'mitsubishi mirage', 'toyota corona mark ii', 'mazda rx3', 'chevrolet
impala']
```

```
[299]: car_names = [name.split()[0] for name in car_names]
print(set(car_names)) # Print unique brand names
```

```
{'buick', 'toyouta', 'saab', 'volkswagen', 'volvo', 'mitsubishi', 'porcshce',
'Nissan', 'porsche', 'mercury', 'renault', 'alfa-romero', 'isuzu', 'mazda',
'jaguar', 'audi', 'subaru', 'maxda', 'vw', 'toyota', 'peugeot', 'dodge',
'plymouth', 'vokswagen', 'chevrolet', 'honda', 'nissan', 'bmw'}
```

```
[ ]:
```

```
[301]: df['CarName'] = df['CarName'].replace({
    'toyouta': 'toyota',
    'porcshce': 'porsche',
    'vokswagen': 'volkswagen',
    'maxda': 'mazda',
    'alfa-romero': 'alfa-romeo',
    'vw': 'volkswagen',
})
```

```
[302]: df
```

```
[302]:
```

	symboling	CarName	fueltype	aspiration	\
74	1	buick regal sport coupe (turbo)	gas	std	
16	0	bmw x5	gas	std	
73	0	buick century special	gas	std	
128	3	porsche boxter	gas	std	
17	0	bmw x3	gas	std	
..	

76	2	mitsubishi mirage	gas	std
150	1	toyota corona mark ii	gas	std
50	1	maxda rx3	gas	std
18	2	chevrolet impala	gas	std
138	2	subaru	gas	std

	doornumber	carbody	drivewheel	enginelocation	wheelbase	carlength	\
74	two	hardtop	rwd	front	112.0	199.2	
16	two	sedan	rwd	front	103.5	193.8	
73	four	sedan	rwd	front	120.9	208.1	
128	two	convertible	rwd	rear	89.5	168.9	
17	four	sedan	rwd	front	110.0	197.0	
..	
76	two	hatchback	fwd	front	93.7	157.3	
150	two	hatchback	fwd	front	95.7	158.7	
50	two	hatchback	fwd	front	93.1	159.1	
18	two	hatchback	fwd	front	88.4	141.1	
138	two	hatchback	fwd	front	93.7	156.9	

	...	enginesize	fuelsystem	boreratio	stroke	compressionratio	\
74	...	304	mpfi	3.80	3.35	8.0	
16	...	209	mpfi	3.62	3.39	8.0	
73	...	308	mpfi	3.80	3.35	8.0	
128	...	194	mpfi	3.74	2.90	9.5	
17	...	209	mpfi	3.62	3.39	8.0	
..	
76	...	92	2bbl	2.97	3.23	9.4	
150	...	92	2bbl	3.05	3.03	9.0	
50	...	91	2bbl	3.03	3.15	9.0	
18	...	61	2bbl	2.91	3.03	9.5	
138	...	97	2bbl	3.62	2.36	9.0	

	horsepower	peakrpm	citympg	highwaympg	price
74	184	4500	14	16	45400.0
16	182	5400	16	22	41315.0
73	184	4500	14	16	40960.0
128	207	5900	17	25	37028.0
17	182	5400	15	20	36880.0
..
76	68	5500	37	41	5389.0
150	62	4800	35	39	5348.0
50	68	5000	30	31	5195.0
18	48	5100	47	53	5151.0
138	69	4900	31	36	5118.0

[205 rows x 25 columns]

```
[303]: # Creating a new feature 'car_stability'
df['car_stability'] = df['wheelbase'] / df['carlength']
df
```

```
[303]:
```

	symboling	CarName	fueltype	aspiration	\
74	1	buick regal sport coupe (turbo)	gas	std	
16	0	bmw x5	gas	std	
73	0	buick century special	gas	std	
128	3	porsche boxter	gas	std	
17	0	bmw x3	gas	std	
..	
76	2	mitsubishi mirage	gas	std	
150	1	toyota corona mark ii	gas	std	
50	1	maxda rx3	gas	std	
18	2	chevrolet impala	gas	std	
138	2	subaru	gas	std	

	doornumber	carbody	drivewheel	engine location	wheelbase	carlength	\
74	two	hardtop	rwd	front	112.0	199.2	
16	two	sedan	rwd	front	103.5	193.8	
73	four	sedan	rwd	front	120.9	208.1	
128	two	convertible	rwd	rear	89.5	168.9	
17	four	sedan	rwd	front	110.0	197.0	
..	
76	two	hatchback	fwd	front	93.7	157.3	
150	two	hatchback	fwd	front	95.7	158.7	
50	two	hatchback	fwd	front	93.1	159.1	
18	two	hatchback	fwd	front	88.4	141.1	
138	two	hatchback	fwd	front	93.7	156.9	

	fuelsystem	boreratio	stroke	compressionratio	horsepower	peakrpm	\
74	mpfi	3.80	3.35	8.0	184	4500	
16	mpfi	3.62	3.39	8.0	182	5400	
73	mpfi	3.80	3.35	8.0	184	4500	
128	mpfi	3.74	2.90	9.5	207	5900	
17	mpfi	3.62	3.39	8.0	182	5400	
..	
76	2bbl	2.97	3.23	9.4	68	5500	
150	2bbl	3.05	3.03	9.0	62	4800	
50	2bbl	3.03	3.15	9.0	68	5000	
18	2bbl	2.91	3.03	9.5	48	5100	
138	2bbl	3.62	2.36	9.0	69	4900	

	citympg	highwaympg	price	car_stability
74	14	16	45400.0	0.562249
16	16	22	41315.0	0.534056
73	14	16	40960.0	0.580971

128	17	25	37028.0	0.529899
17	15	20	36880.0	0.558376
..
76	37	41	5389.0	0.595677
150	35	39	5348.0	0.603025
50	30	31	5195.0	0.585167
18	47	53	5151.0	0.626506
138	31	36	5118.0	0.597196

[205 rows x 26 columns]

```
[311]: # Dropping highly correlated features
drop_features = ['carlength', 'carwidth', 'curbweight', 'wheelbase',
↳ 'highwaympg']
df = df.drop(columns=drop_features, axis=1)
```

```
[313]: # Encoding categorical variables using dummy encoding
df = pd.get_dummies(df, drop_first=True)
```

```
[318]: df.columns = df.columns.str.replace(' ', '_')
df
```

```
[318]:
```

	symboling	carheight	enginesize	boreratio	stroke	compressionratio	\
74	1	55.4	304	3.80	3.35	8.0	
16	0	53.7	209	3.62	3.39	8.0	
73	0	56.7	308	3.80	3.35	8.0	
128	3	51.6	194	3.74	2.90	9.5	
17	0	56.3	209	3.62	3.39	8.0	
..	
76	2	50.8	92	2.97	3.23	9.4	
150	1	54.5	92	3.05	3.03	9.0	
50	1	54.1	91	3.03	3.15	9.0	
18	2	53.2	61	2.91	3.03	9.5	
138	2	53.7	97	3.62	2.36	9.0	

	horsepower	peakrpm	citympg	price	...	cylindernumber_three	\
74	184	4500	14	45400.0	...	False	
16	182	5400	16	41315.0	...	False	
73	184	4500	14	40960.0	...	False	
128	207	5900	17	37028.0	...	False	
17	182	5400	15	36880.0	...	False	
..	
76	68	5500	37	5389.0	...	False	
150	62	4800	35	5348.0	...	False	
50	68	5000	30	5195.0	...	False	
18	48	5100	47	5151.0	...	True	
138	69	4900	31	5118.0	...	False	

	cylindernumber_twelve	cylindernumber_two	fuelsystem_2bbl	\
74	False	False	False	
16	False	False	False	
73	False	False	False	
128	False	False	False	
17	False	False	False	
..	
76	False	False	True	
150	False	False	True	
50	False	False	True	
18	False	False	True	
138	False	False	True	


	fuelsystem_4bbl	fuelsystem_idi	fuelsystem_mfi	fuelsystem_mphi	\
74	False	False	False	True	
16	False	False	False	True	
73	False	False	False	True	
128	False	False	False	True	
17	False	False	False	True	
..	
76	False	False	False	False	
150	False	False	False	False	
50	False	False	False	False	
18	False	False	False	False	
138	False	False	False	False	

	fuelsystem_spdi	fuelsystem_spfi
74	False	False
16	False	False
73	False	False
128	False	False
17	False	False
..
76	False	False
150	False	False
50	False	False
18	False	False
138	False	False

[205 rows x 186 columns]

```
[320]: # Data Analysis - Histogram & Bar Charts
plt.figure(figsize=(10, 6))
df['price'].hist(bins=30, edgecolor='black')
plt.xlabel('Price')
plt.ylabel('Frequency')
```

```
plt.title('Distribution of Car Prices')  
plt.show()
```



```
[324]: # Generate Insights and save the report  
df.describe().to_csv('data_analysis_report.csv')
```

```
[ ]:
```

```
[ ]:
```